

# Light virtualization (containers)

Davide Rossi

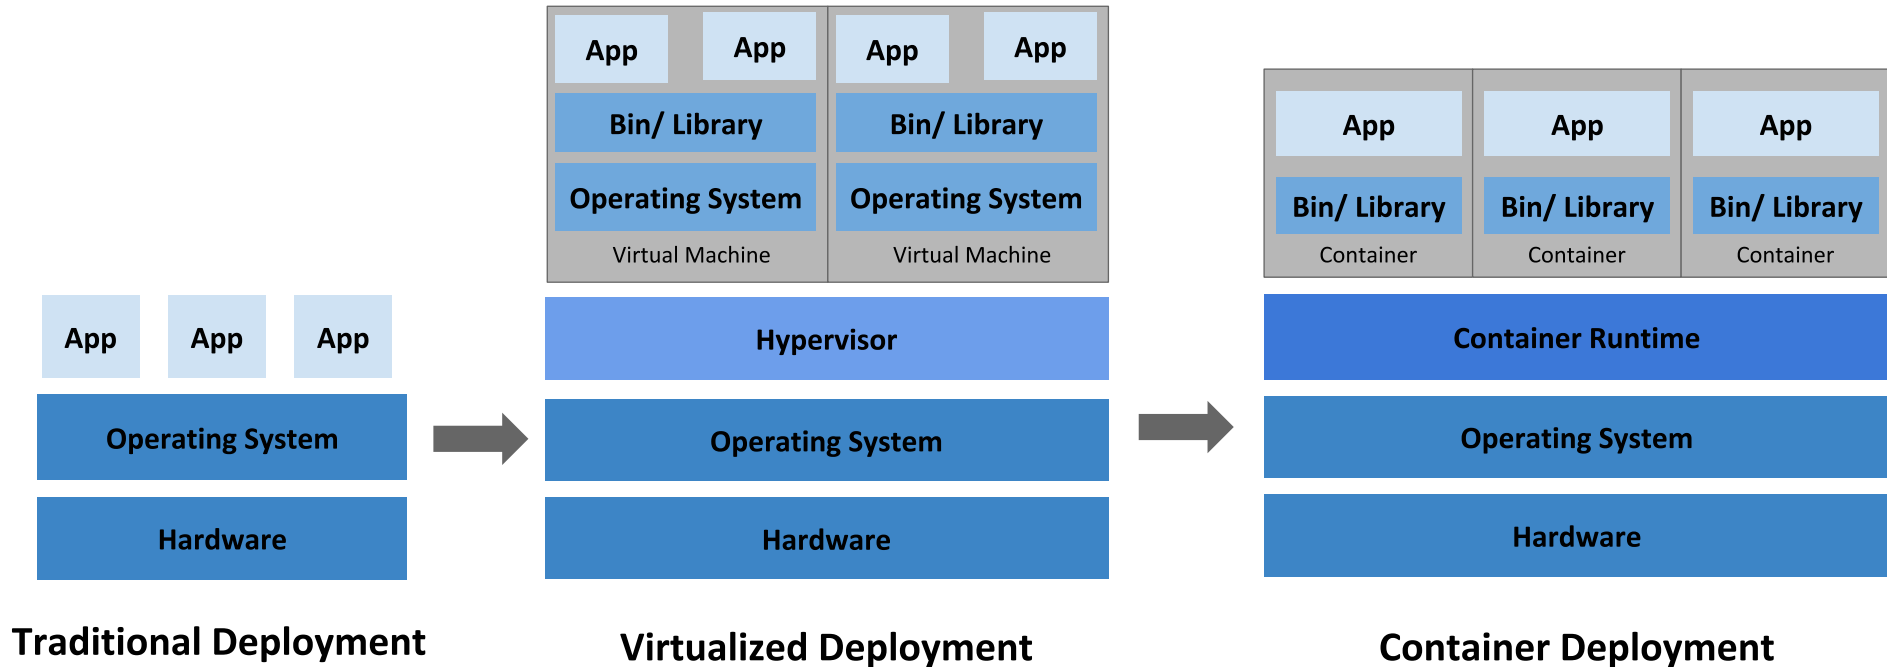
Dipartimento di Informatica – Scienze e Ingegneria  
Università di Bologna



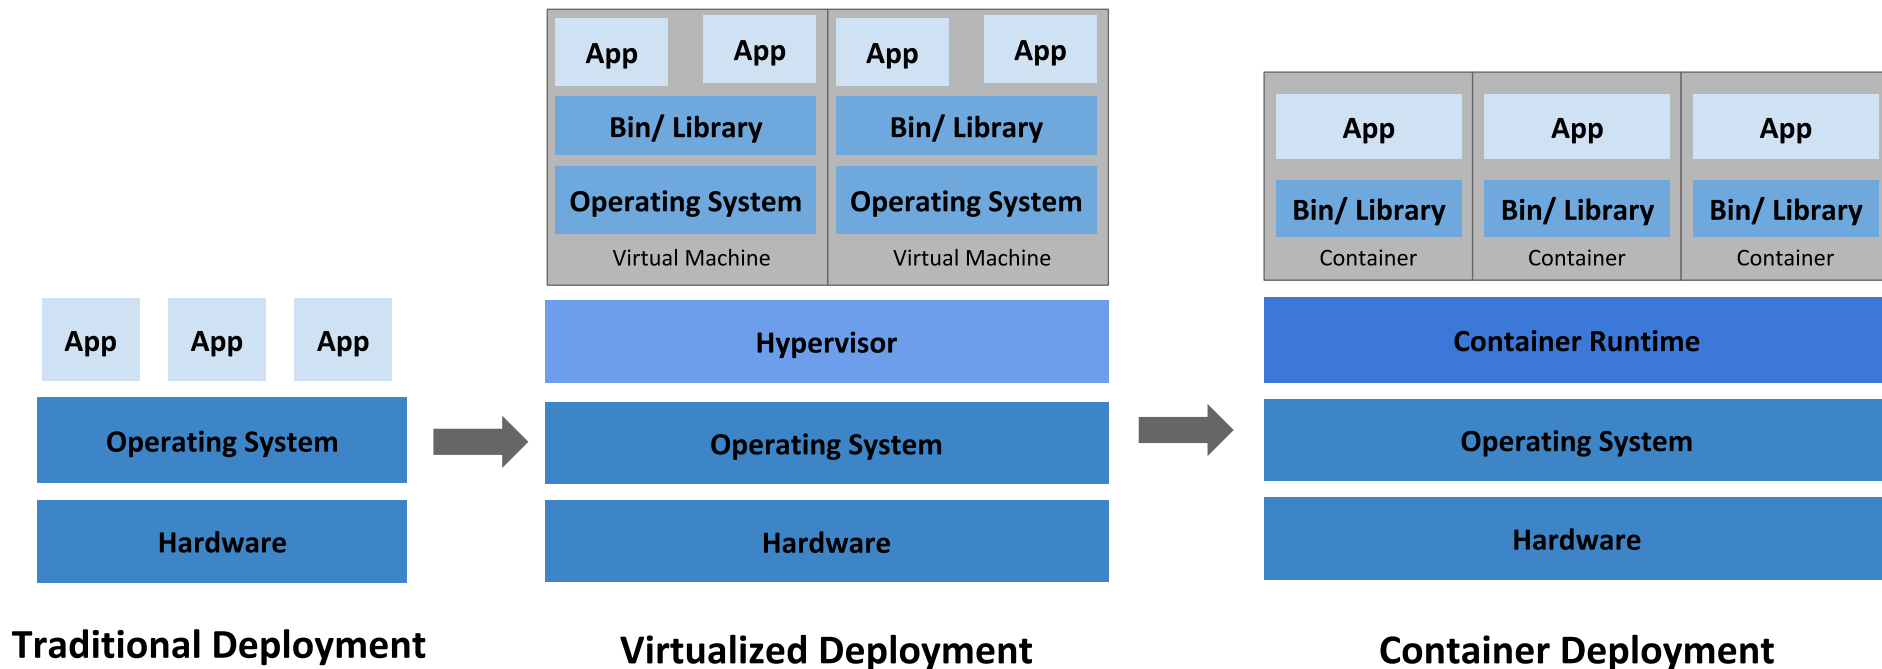
# What is a container

- A container is an isolated partition of the OS
- Containers are directly supported by the OS (that is: Docker is just a management tool)

# Containers and *light* virtualization



# Containers and *light* virtualization



Notice that on Windows containers are really *light* VMs, i.e. they run on top of a hypervisor.

# Containers on Linux

Linux manages containers using

- namespaces
  - create partitions for OS managed abstractions
- cgroups
  - limit processes access to resources

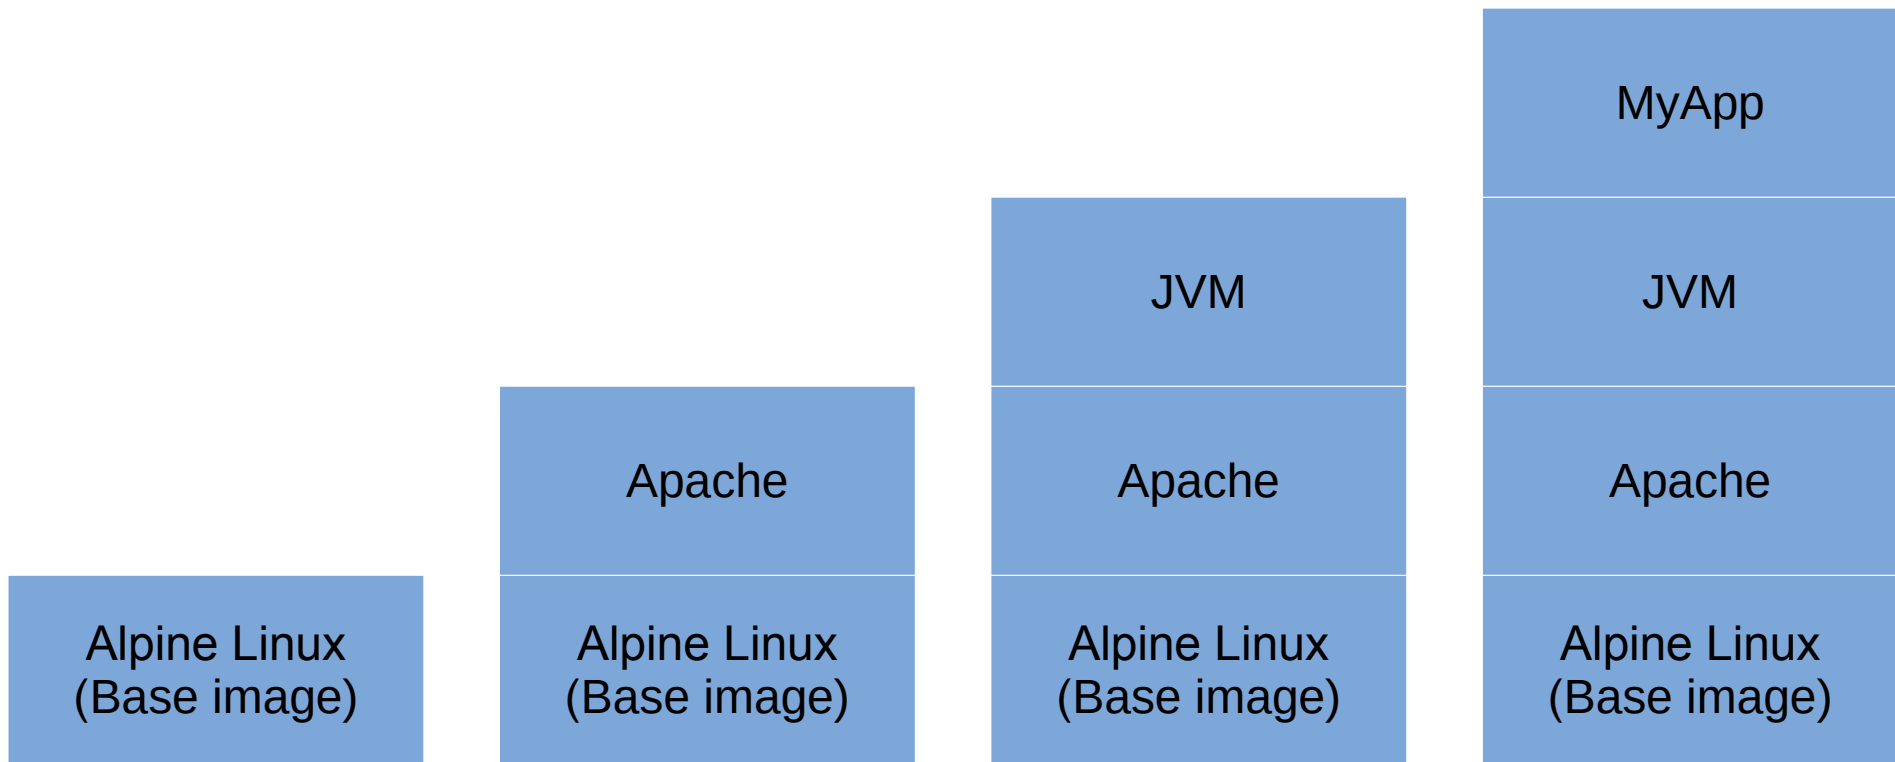
# Filesystem

- A container usually has its own FS
- While this could just be a chrooted partition of the host FS, it is usually a loopback/*layered* filesystem

# Container images

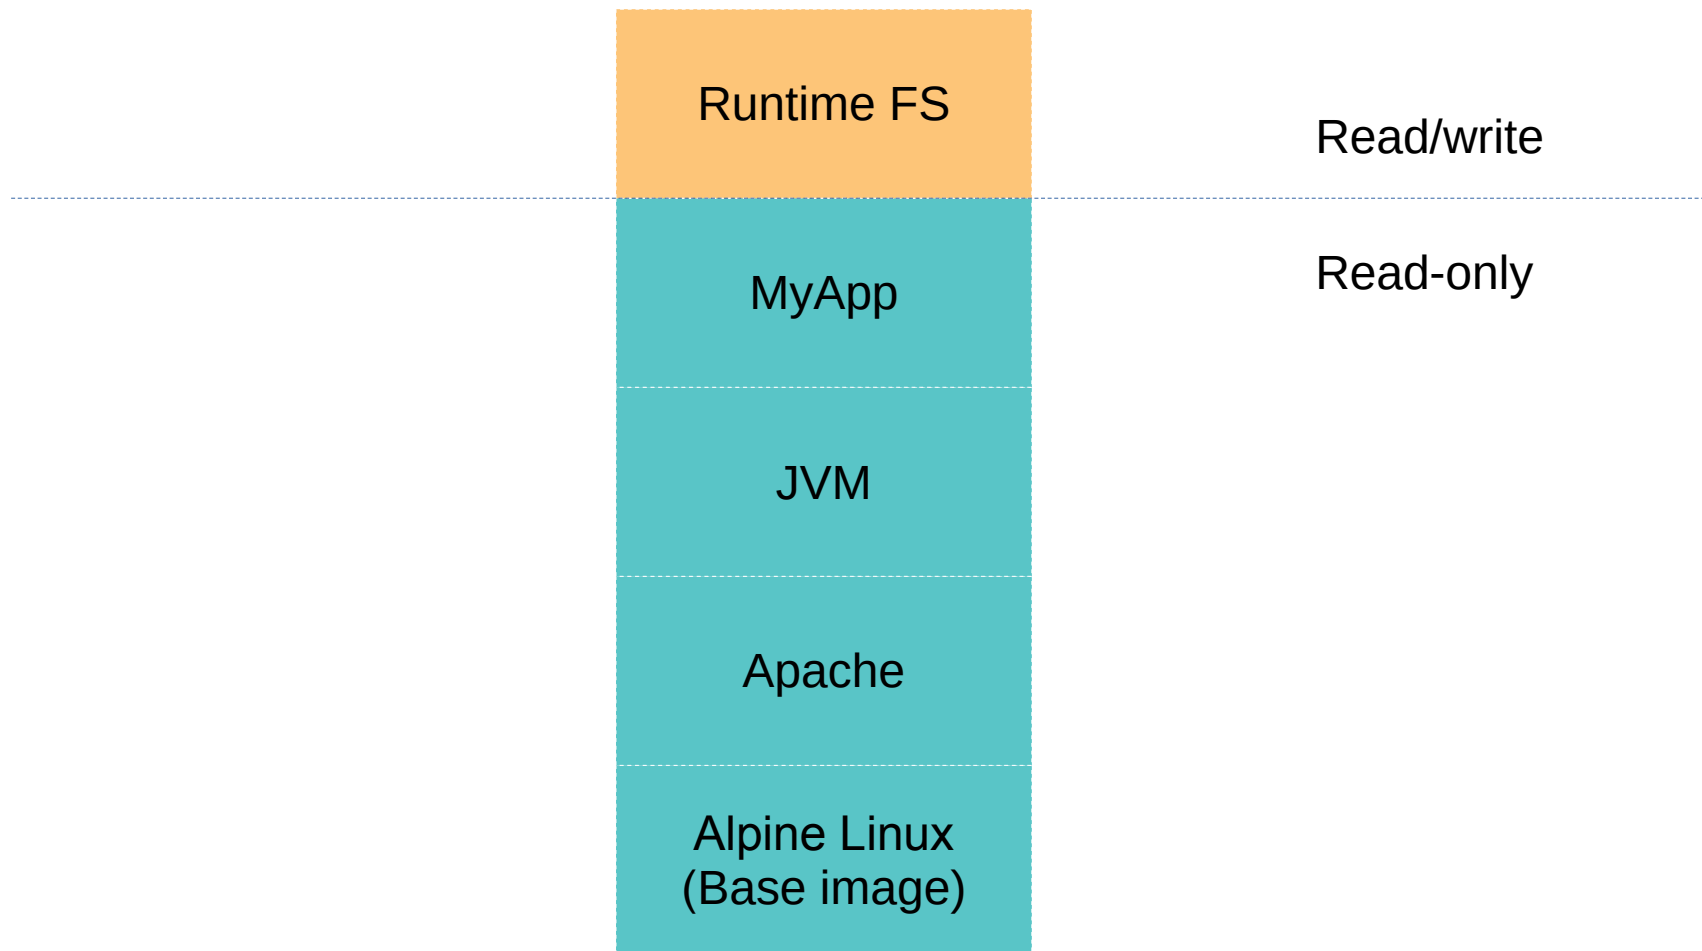
- Bundles containing all that is needed to setup a container
  - FS content
  - Initial process
  - Misc metadata & configuration bits

# Layered FS images

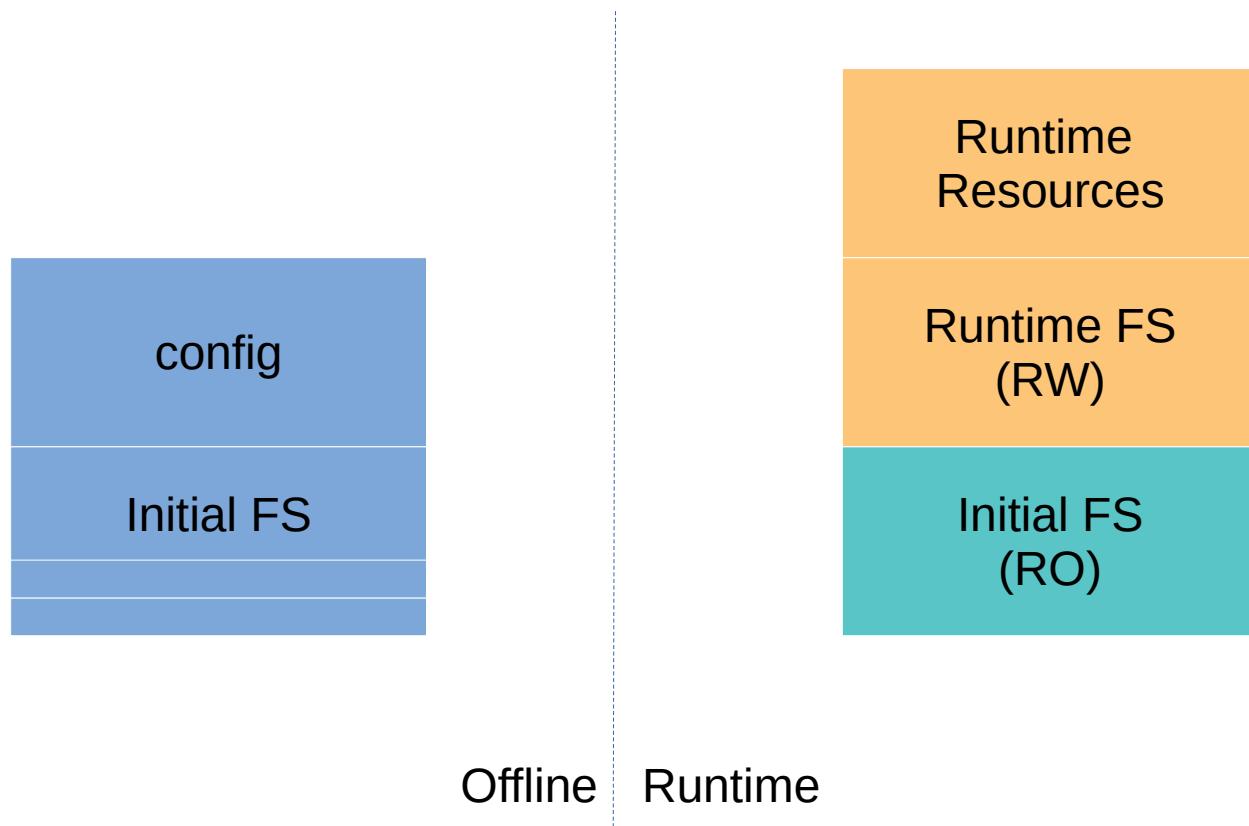




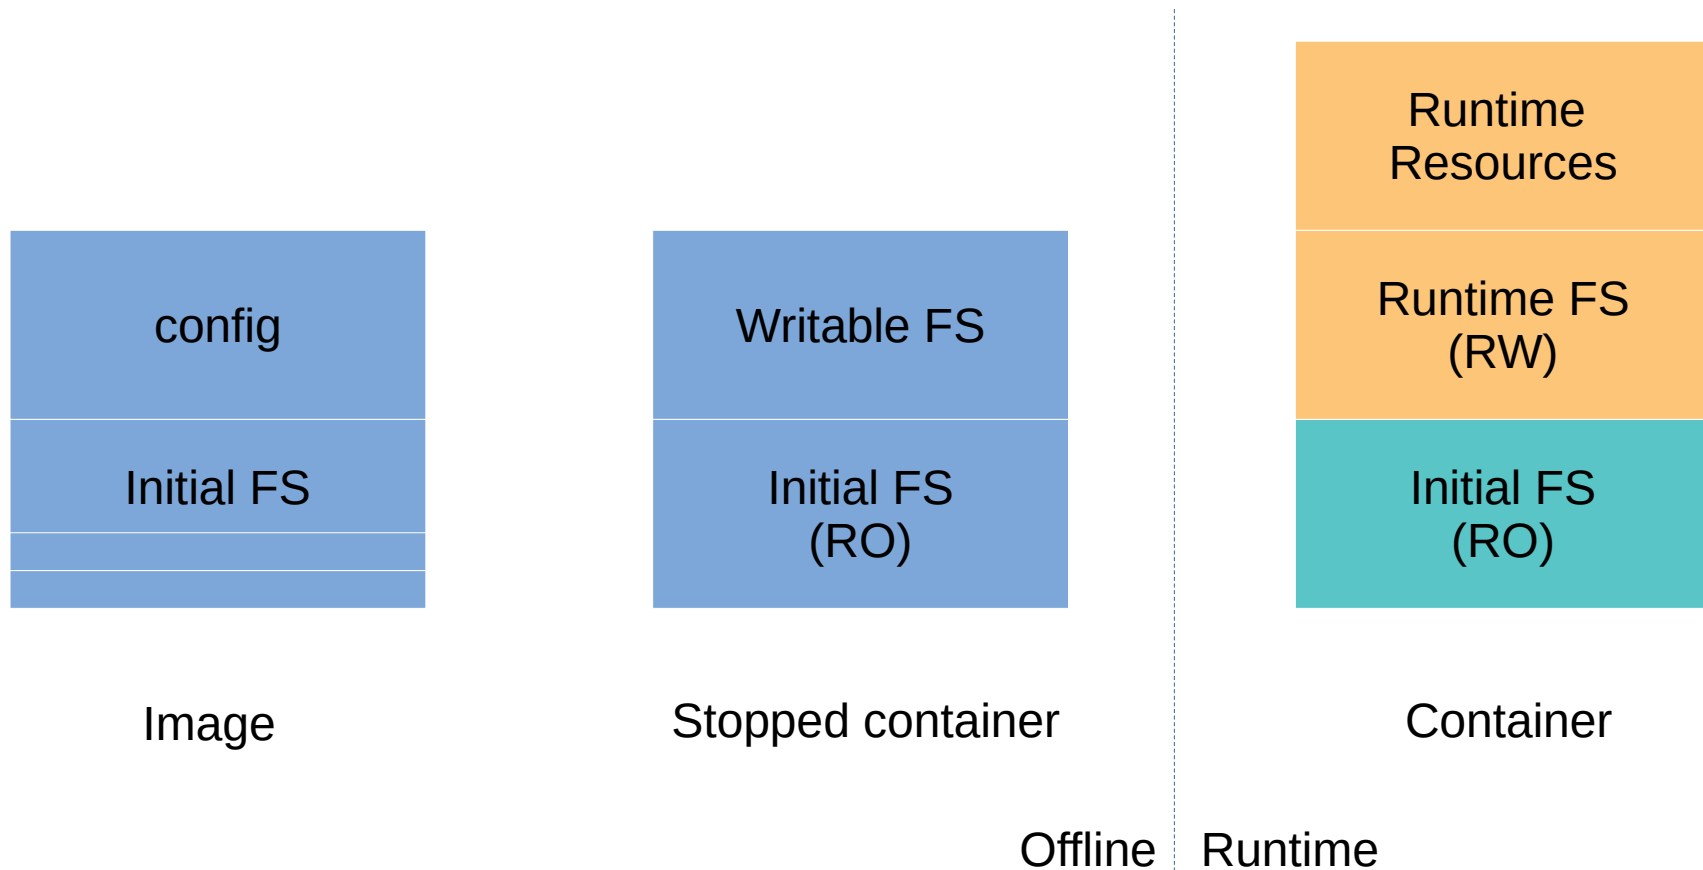
# Runtime FS



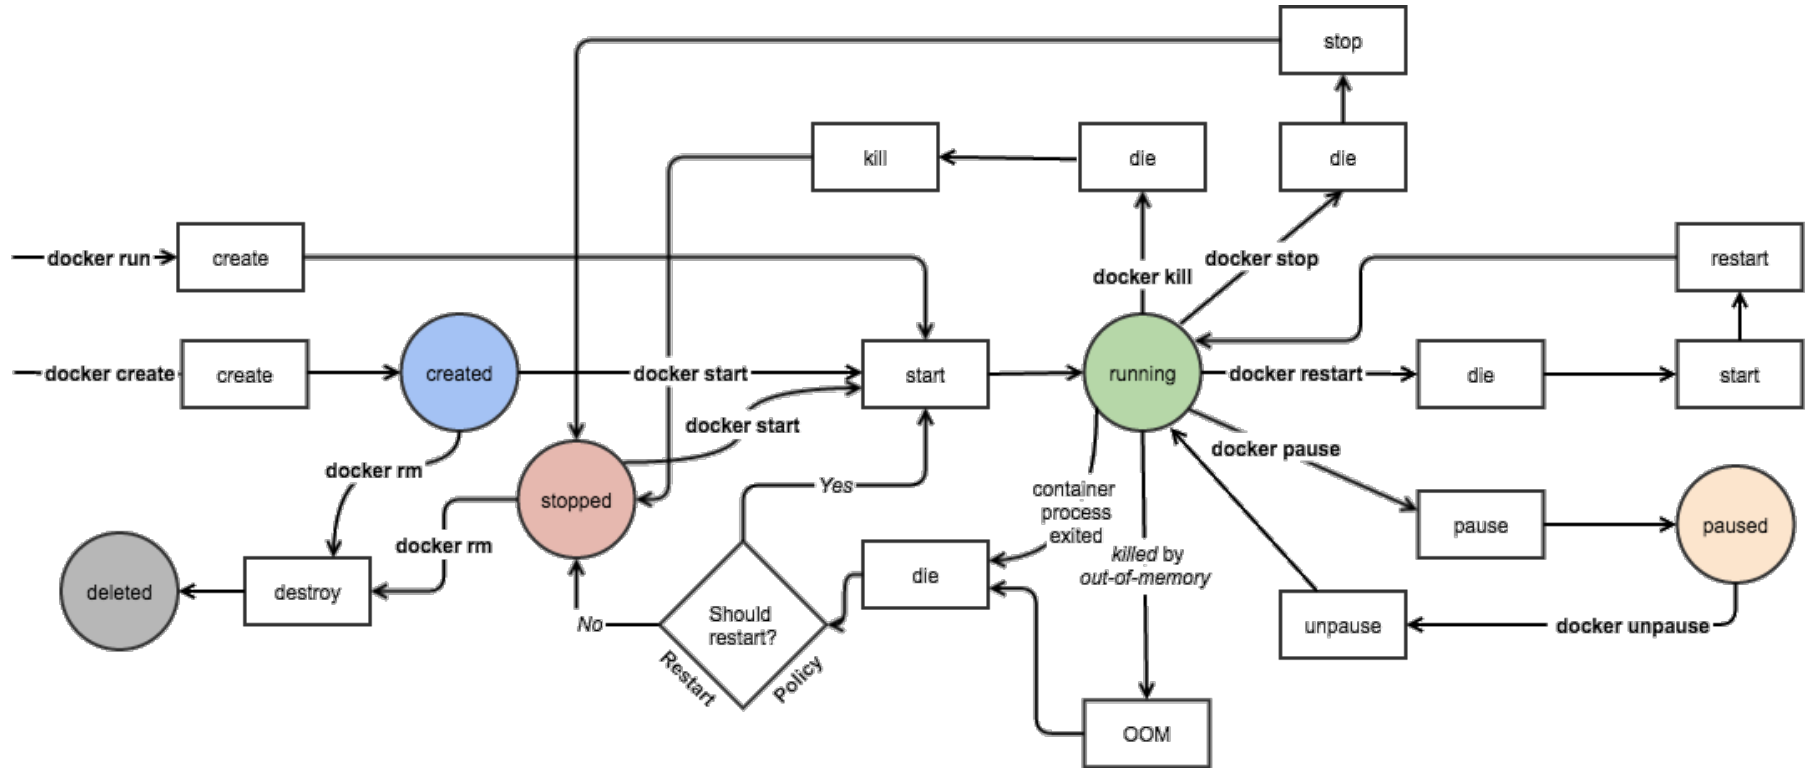
# Images vs containers



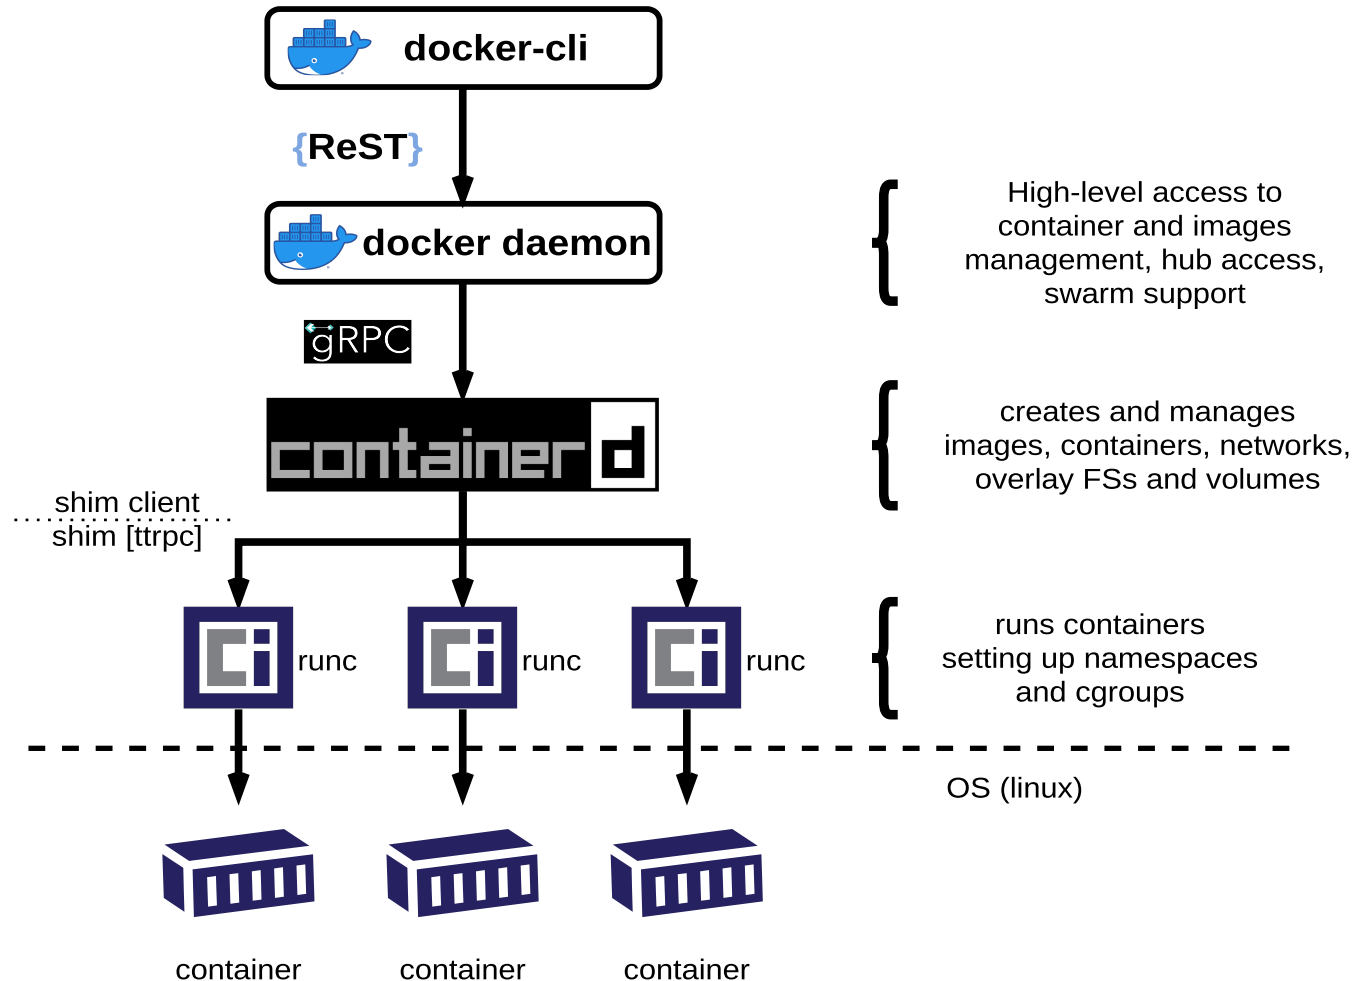
# The Docker view



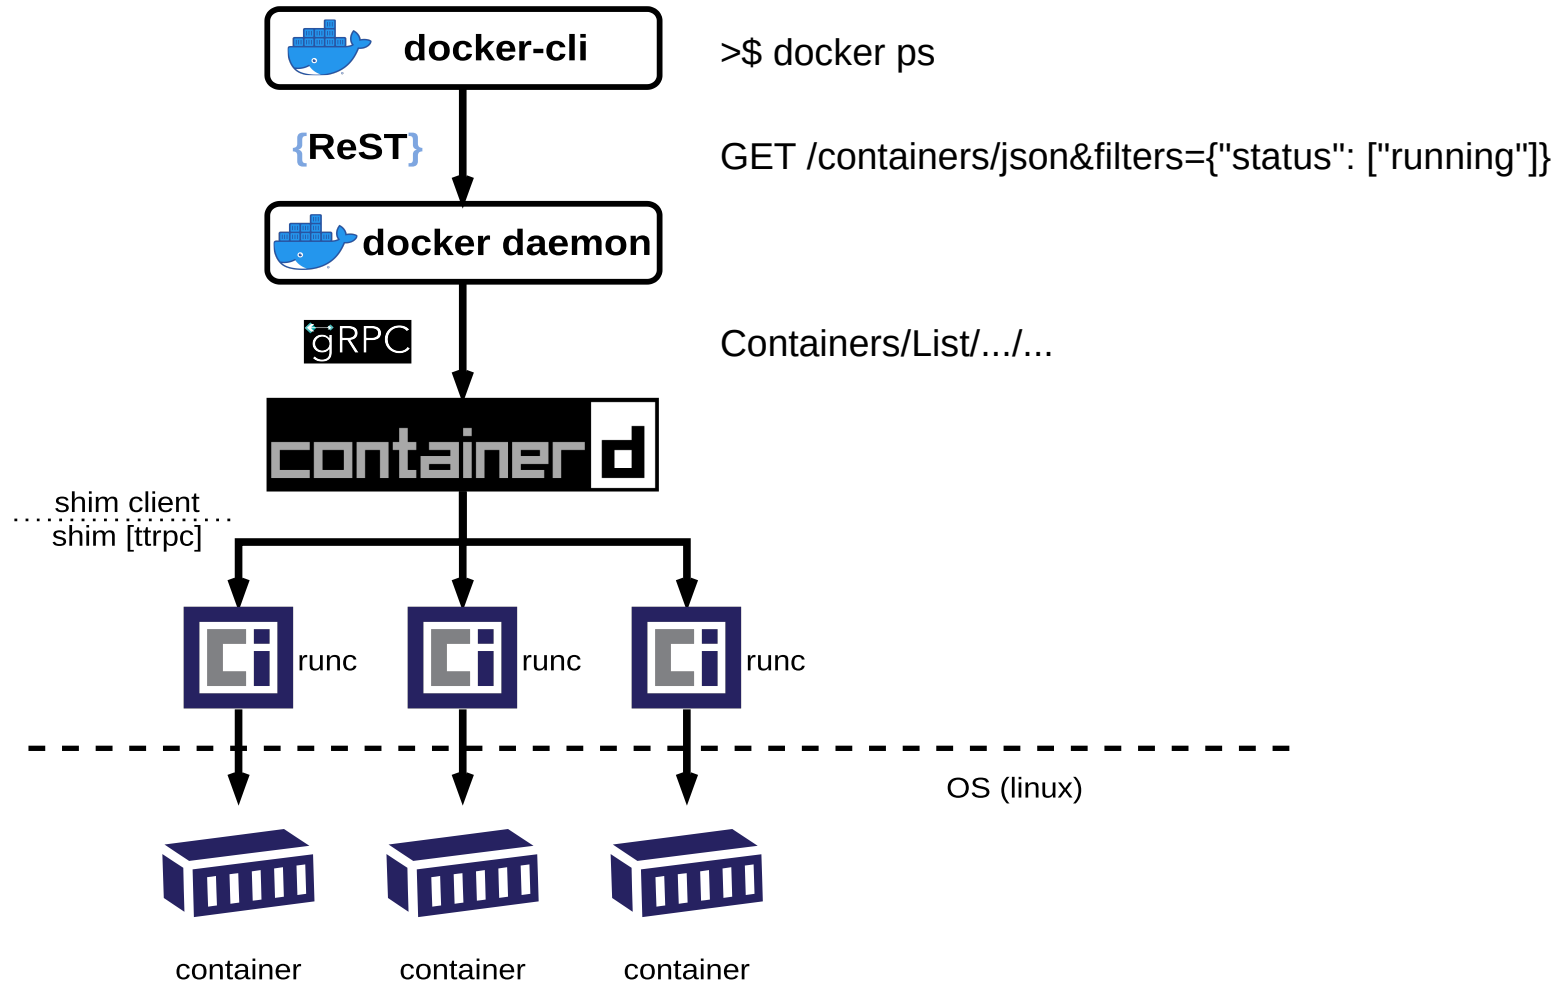
# Specifically...



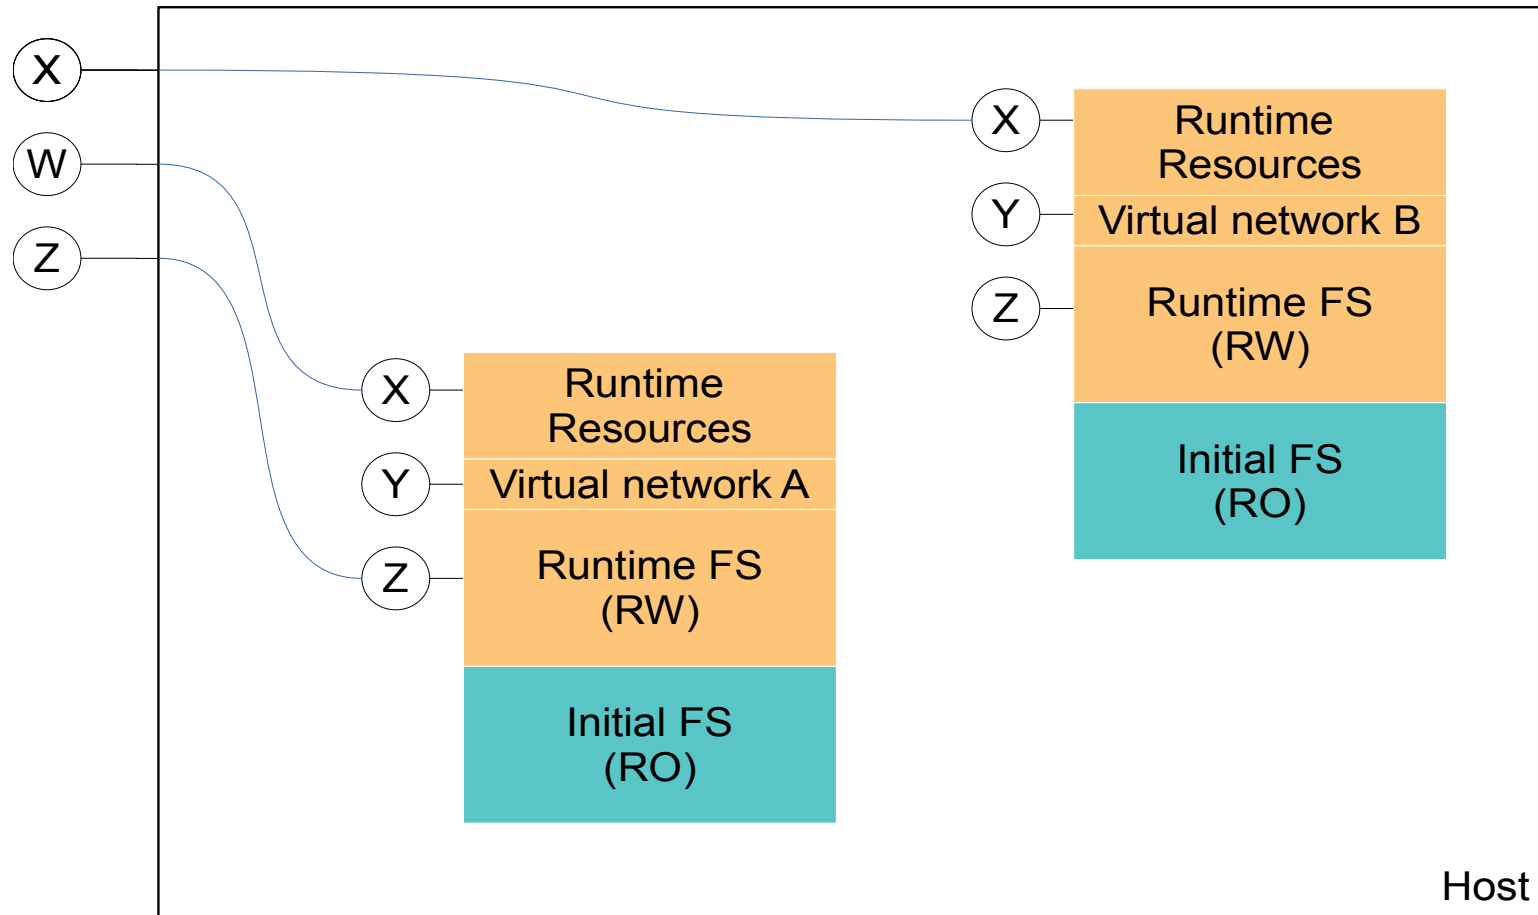
# The Docker stack



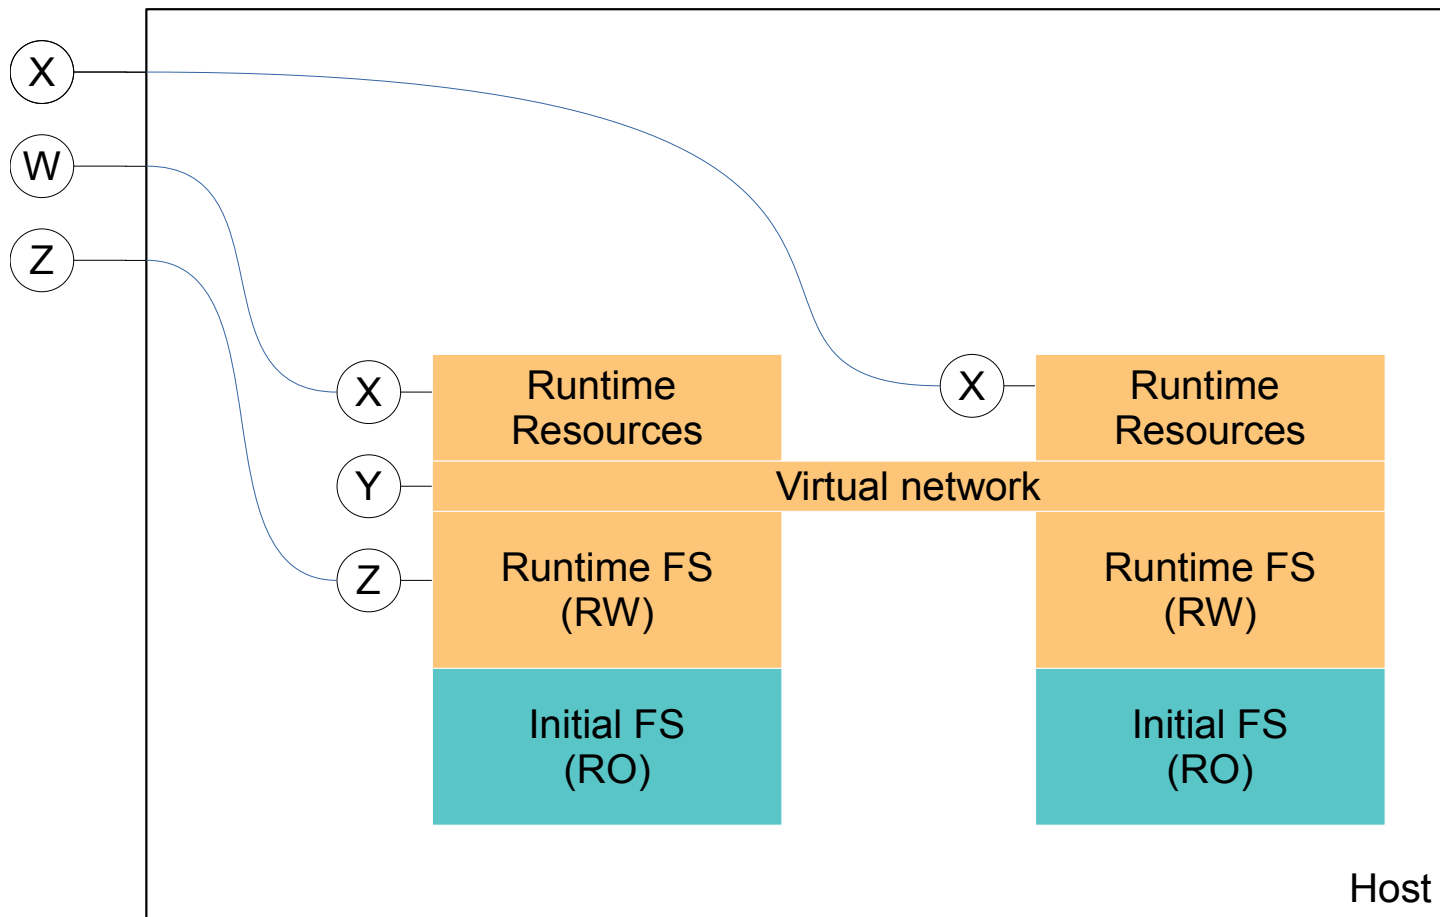
# The Docker stack



# Networking



# Intra containers





# Dockerfile

- Note: using docker and a dockerfile is one of the possible ways to create a container image compliant with the Open Container Initiative (OCI) image specification
- A dockerfile uses a DSL (the Dockerfile syntax) that is parsed by Docker (or other tools) to create container images

# Dockerfile structure

```
1 #syntax=docker/dockerfile:1
2 FROM golang:1.21-alpine
3 WORKDIR /src
4 COPY . .
5 RUN go mod download
6 RUN go build -o /bin/client ./cmd/client
7 RUN go build -o /bin/server ./cmd/server
8 ENTRYPOINT [ "/bin/server" ]
```

# Multi-stage builds

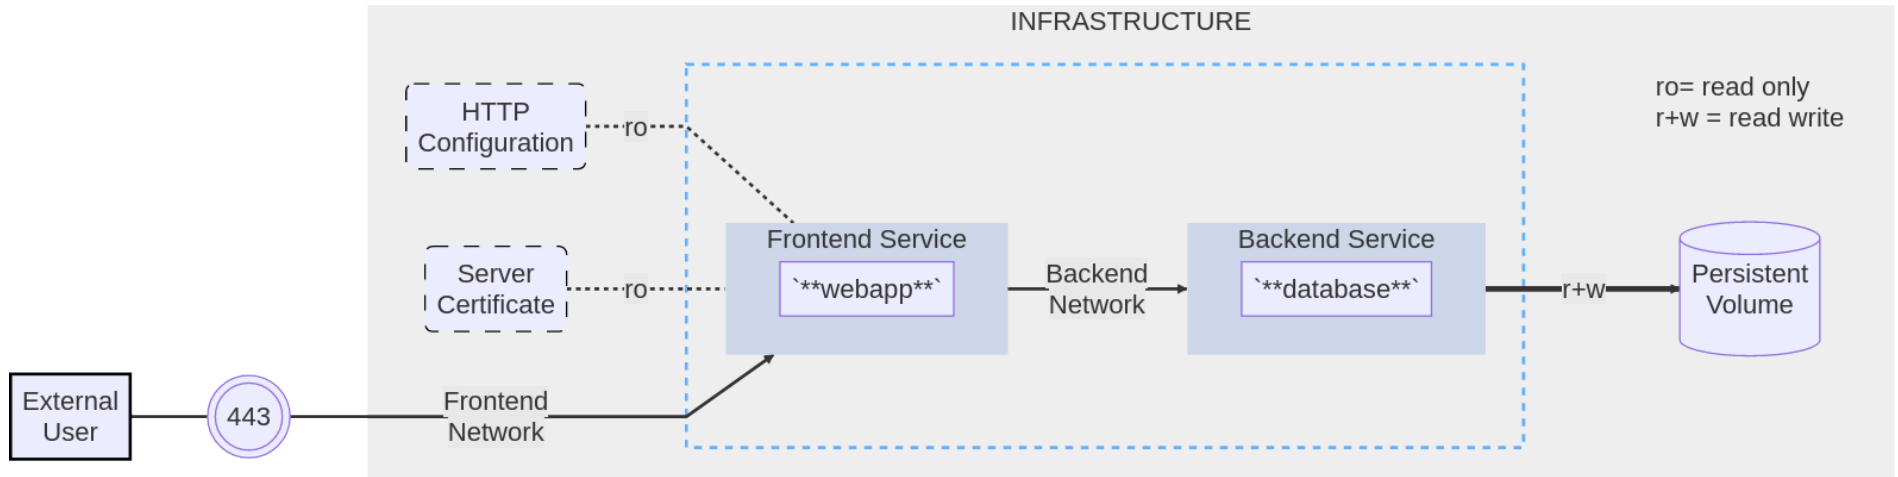
```
1 #syntax=docker/dockerfile:1
2 FROM golang:1.21-alpine as build
3 WORKDIR /src
4 COPY . .
5 RUN go mod download
6 RUN go build -o /bin/client ./cmd/client
7 RUN go build -o /bin/server ./cmd/server
8
9 FROM scratch
10 COPY --from=build /bin/client /bin/server /bin/
11 ENTRYPOINT [ "/bin/server" ]
```

# Compose

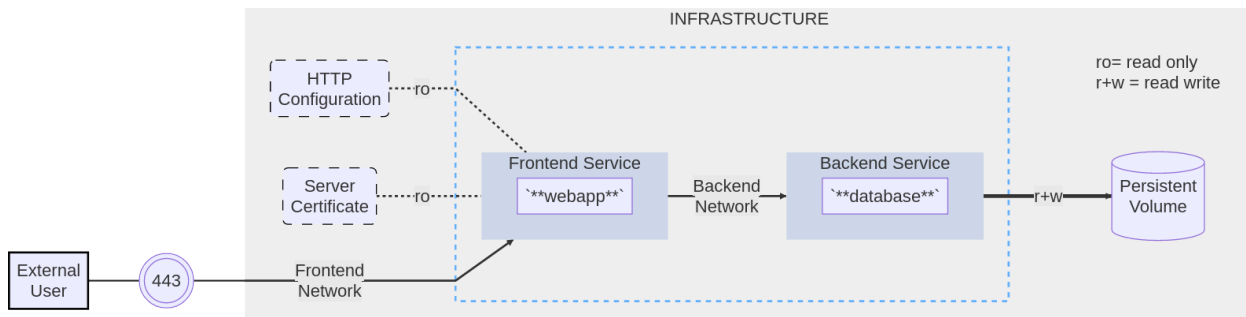
A standardized specification for container-based applications (Docker compose is the reference implementation).

It allows to declaratively describe how to set up and execute an application composed by multiple cooperating containers.

# Compose example



# Compose example



```
services:
  frontend:
    image: awesome/webapp
    ports:
      - "443:8043"
    networks:
      - front-tier
      - back-tier
    configs:
      - httpd-config
    secrets:
      - server-certificate
```

```
backend:
  image: awesome/database
  volumes:
    - db-data:/etc/data
  networks:
    - back-tier
```

```
volumes:
  db-data:
    driver: flocker
    driver_opts:
      size: "10GiB"
```

```
configs:
  httpd-config:
    external: true
```

```
secrets:
  server-certificate:
    external: true
```

```
networks:
  # The presence of these objects is sufficient to define them
  front-tier: {}
  back-tier: {}
```

# Kubernetes (K8s)

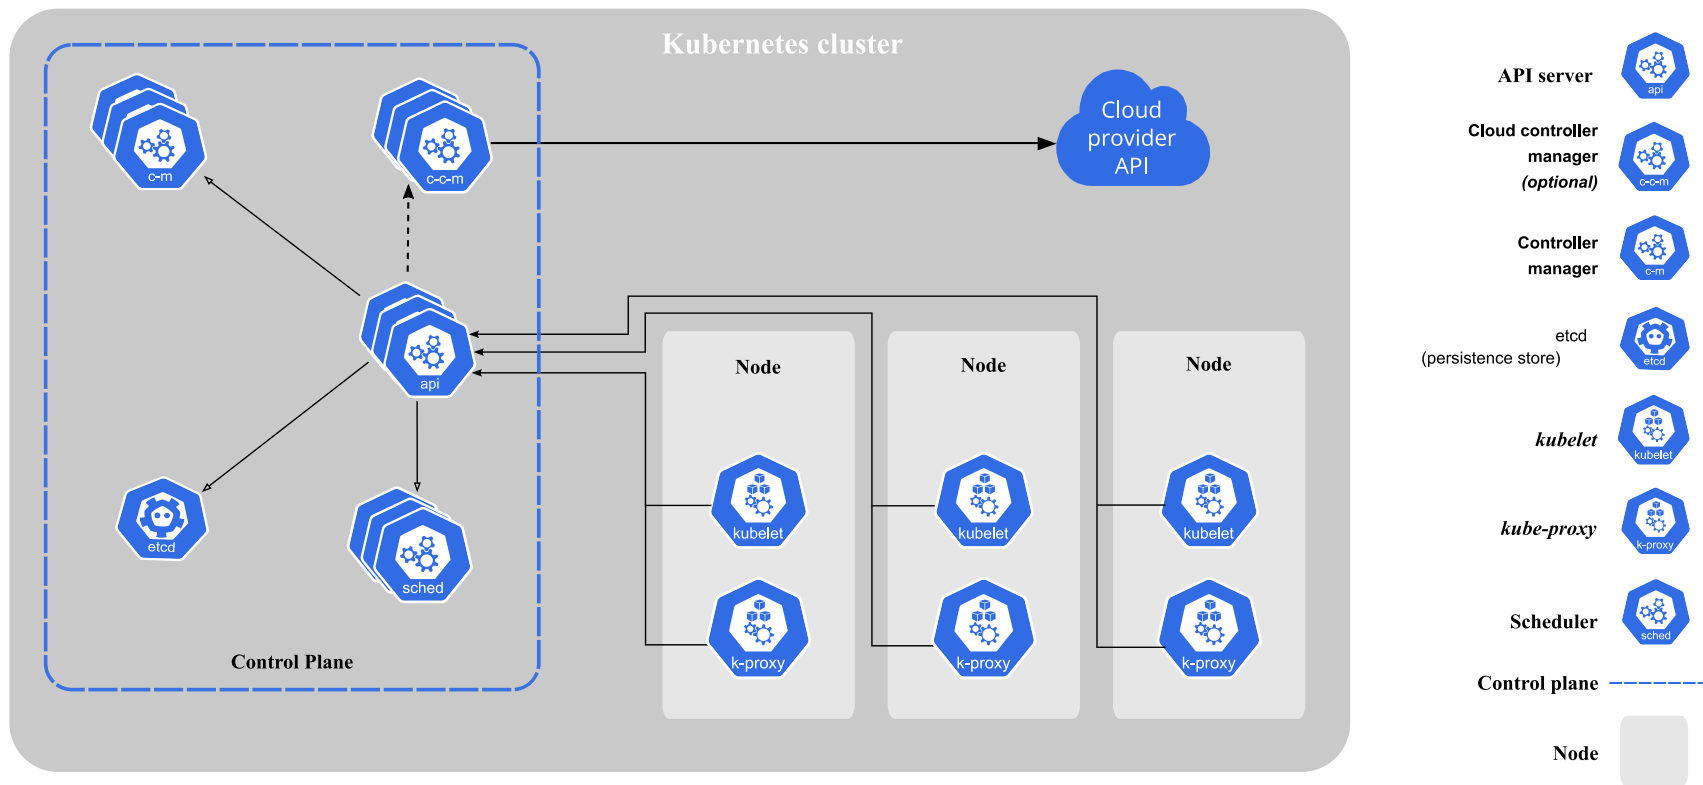
- Kubernetes allocates containers in multiple *nodes*

# Kubernetes (K8s)

- Kubernetes allocates *containerized applications* in a *cluster* composed by multiple *nodes*
- Providing support for automating deployment, scaling, networking and management



# K8s components



# K8s

- Kubernetes manages the state of the system as a set of persistent *objects*
- Kind of objects in K8s:  
Pod, Deployment, ReplicaSet, StatefulSet, DaemonSet, PersistentVolume, Service, Ingress, Namespace, ConfigMap, Secret, Job

# K8s general concepts

- Object management techniques
  - Imperative (commands/configurations)
  - Declarative configurations – a desired state is declared, K8s' machinery works to reach that state
- A simple matching mechanism between *labels* (key-value pairs) and *selectors* is used to associate objects

# Pod objects

- Pods are ephemeral deployment units
- A pod is a group of one or more containers, with shared storage and network resources
- Pods are always co-located and co-scheduled
- *Think of them as a compose with **less** isolation*

# Isolation objects

- Namespace
  - Isolate partition within a cluster

# Workload objects

- ReplicaSet
  - Stable set of replica pods
- Deployment
  - Declarative updates for Pods and ReplicaSets
- StatefulSet
  - Manages replicas with identities

# Workload objects

- DaemonSet
  - Manages pod to be allocated to nodes
- Job
  - Manages (retriable) pods executions

# Storage objects

- PersistentVolume
  - Abstract storage resources (to be matched by corresponding persistentVolumeClaims)



# Configuration objects

- ConfigMap
  - Key-value storage that can be consumed by pods
- Secret
  - Key-value storage intended to hold confidential data

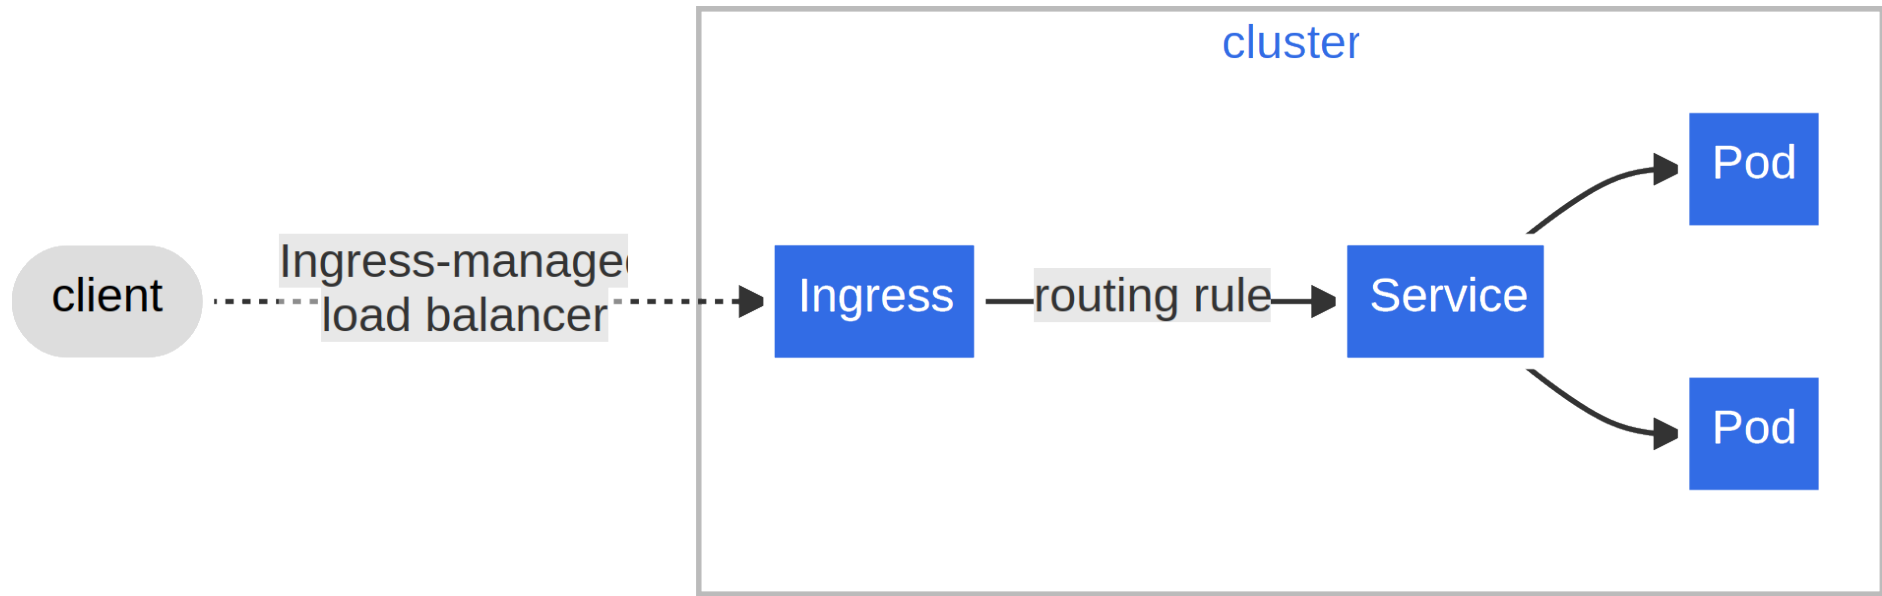
# Access objects

- Service
  - Exposes a network application (running in pods in the cluster)  
Services are registered in the cluster's internal DNS
- Ingress
  - Manages external access to services in a cluster

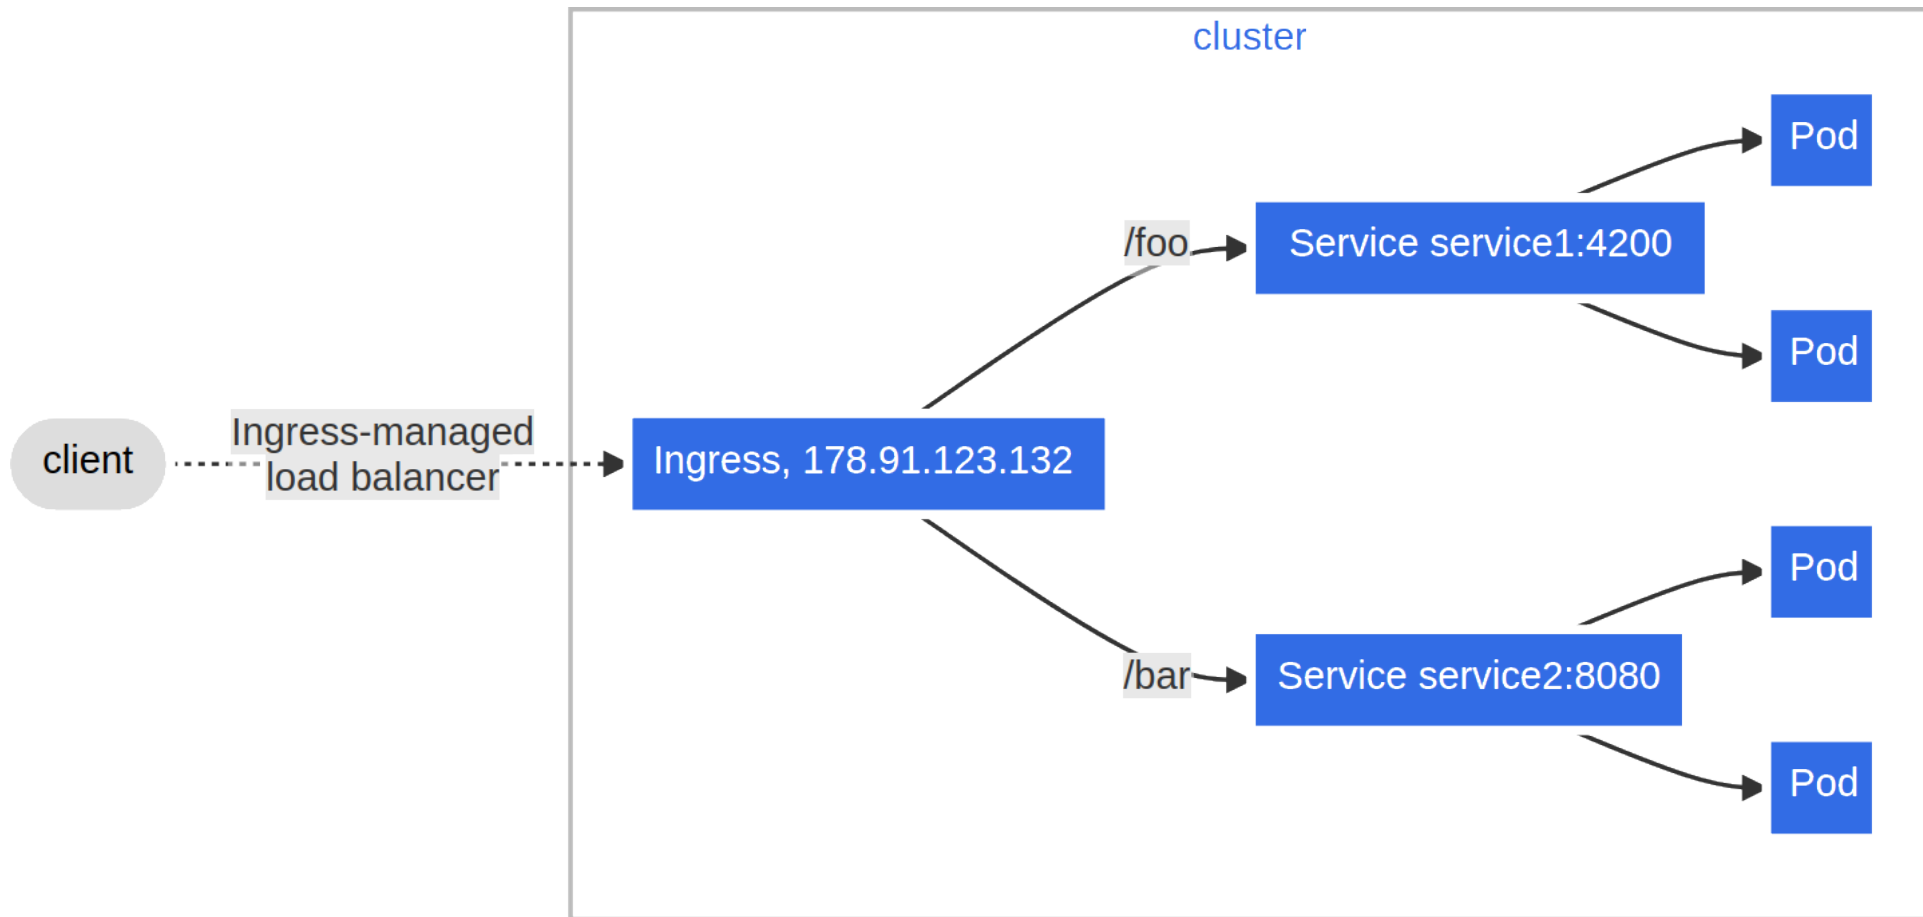
# Service types

- ClusterIP
  - Allocates a cluster-wide internal virtual IP address on all nodes. Traffic to this IP is redirected to the pod's endpoints via OS-provided mechanisms
- NodePort
  - Exposes the ClusterIP on all nodes (on the same port). Incoming traffic is subject to NAT
- LoadBalancer
  - Configures an external load balancer to access NodePorts

# Ingress



# Ingress



# Ingress

