

# Services and communications

Davide Rossi

Dipartimento di Informatica – Scienze e Ingegneria  
Università di Bologna



# Topics

- Communication styles
- REST
- Binary messaging

# Communication styles

- Explicit addressing
  - Request/response (a.k.a. RPC)
    - Sync: the response uses the request's channel
    - Async: the response uses a different channel (the caller must expose a response endpoint)
  - Streaming
- Implicit addressing
  - Queues
  - Publish/subscribe channels (topics)
  - Logs

# Communication styles

- Explicit addressing
    - Request/response (a.k.a. RPC)
      - Sync: the response uses the request's channel
      - Async: the response uses a different channel (the caller must expose a response endpoint)
    - Streaming
  - Implicit addressing
    - Queues
    - Publish/subscribe channels (topics)
    - Logs
- } Message brokers
- } Distributed event logs / streaming platforms

# Communication styles

- Explicit addressing
    - Request/response (a.k.a. RPC)
      - Sync: the response uses the request's channel
      - Async: the response uses a different channel (the caller must expose a response endpoint)
    - Streaming
  - Implicit addressing
    - Queues
    - Publish/subscribe channels (topics)
    - Logs
- } Event-driven architectural style


# Event-driven architectural style

- The main components are event **producers** and event **consumers**, which are decoupled (producers are unaware of consumers and vice versa)
- **Pub/sub model**: consumers register (subscribe) with a message infrastructure. Generated (published) events are dispatched to registered consumers. Subscribers only receive messages published after their subscription
- **Event streaming model**: generated events are persisted in a **log** where they are stored in order (typically within a *partition*). Clients can process the events by retrieving them from the stream and are responsible for positioning (which means that they can replay events generated at any time)


# Message brokers




Queue




Work queue  
(competing consumer group)




Publish/subscribe



Routing (queue filter)



Topics (pattern matching filter)



RPC

# REST

REST is an architectural style.

The largest known implementation of a system conforming to the REST architectural style is the World Wide Web.

REST exemplifies how the Web's architecture emerged by characterizing and constraining the macro-interactions of the four components of the Web, namely origin servers, gateways, proxies and clients, without imposing limitations on the individual participants. As such, REST essentially governs the proper behavior of participants.



# Architectural Style

A set of design rules that identify the kinds of components and connectors that may be used to compose a system or subsystem, together with local or global **constraints** on the way the composition is done.


[Shaw & Clements, 1996]

# REST constraints

The REST architectural style describes the following six constraints applied to the architecture, while leaving the implementation of the individual components free to design:

- Client–server
- Stateless
- Cacheable
- Uniform interface
- Layered system
- Code on demand (optional)

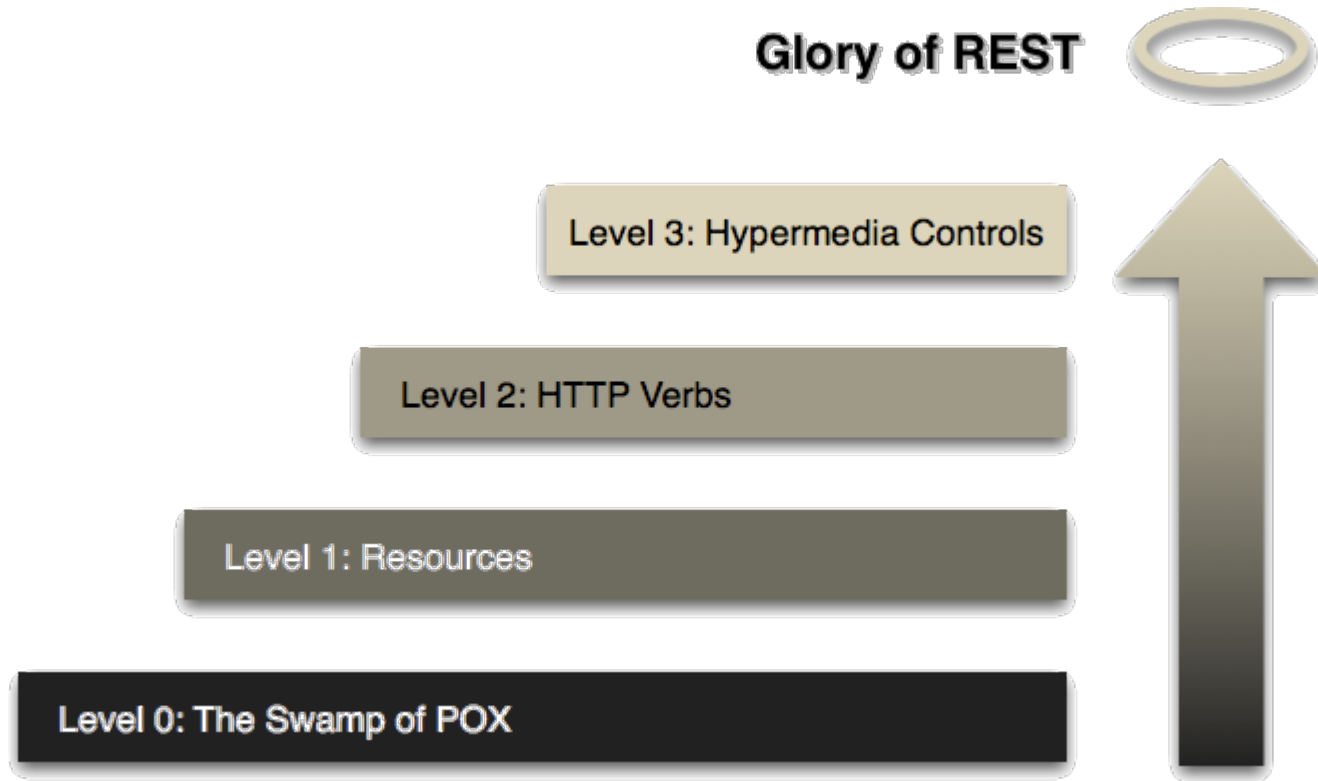
# (Simplified) Web architecture



# REST and the Web

REST	Web
Client–server	HTTP
Stateless	HTTP
Cacheable	HTTP/Web architecture
Uniform interface	HTTP/URI/HATEOAS/ content negotiation
Layered system	Web architecture
Code on demand	Javascript (?)

# Richardson Maturity Model



# HATEOAS

- Hypermedia as the engine of application state (HATEOAS): resource representations contain links to other resources.
- Links make interconnected resources navigable. Identifying new resources is service-specific. RESTful applications navigate instead of calling. Representations contain information about possible traversals. The application navigates to the next resource depending on link semantics. Navigation can be delegated since all links use identifiers.

# REST and CRUD

Create	POST	Create a resource
Read	GET	Retrieve the current state of a resource
Update	PUT (PATCH)	Initialize or update the state of a resource
Delete	DELETE	Remove a resource

# REST factory pattern

If we want to create a resource we could use PUT but we have to make sure that the new URI we provide is not already allocated and that the server is not planning on using it in the future.

The POST factory pattern can help in this context:

→ POST /resource

← 201 Created

Location: /resource/<newid>



# Example: Doodle

	GET	PUT	POST	DELETE
<code>/poll</code>	V	X	V	X
<code>/poll/{pid}</code>	V	V	X	V
<code>/poll/{pid}/vote</code>	V	X	V	X
<code>/poll/{pid}/vote/{vid}</code>	V	V	X	?

# REST interface

- Within RMML2 resource paths are contracts
- If a server moves a resource to another path the clients break
- Thus, **paths are interfaces**
- Yet there is no standard notation for REST interfaces
  - Usual approach: read the docs. Well, docs are not machine processable!

# Describing REST API

- WADL (W3C)
- RAML
- OAS - OpenAPI Specification (formerly Swagger)
- API Blueprint

Machine



Human

All formats are machine processable and can be used to generate documentation and development artifacts (skel/proxy, ...)

# OAS Petstore

```
1 openapi: "3.0.0"
2 info:
3   version: 1.0.0
4   title: Swagger Petstore
5   license:
6     name: MIT
7 servers:
8   - url: http://petstore.swagger.io/v1
9 paths:
10  /pets:
11    get:
12      summary: List all pets
13      operationId: listPets
14      tags:
15        - pets
16      parameters:
17        - name: limit
18          in: query
19          description: How many items to return at one time (max 100)
20          required: false
21          schema:
22            type: integer
23            format: int32
24      responses:
25        '200':
26          description: A paged array of pets
27          headers:
28            x-next:
29              description: A link to the next page of responses
30              schema:
31                type: string
32          content:
33            application/json:
34              schema:
35                $ref: "#/components/schemas/Pets"
36        default:
37          description: unexpected error
38          content:
39            application/json:
40              schema:
41                $ref: "#/components/schemas/Error"
42    post:
43      summary: Create a pet
44      operationId: createPets
45      tags:
46        - pets
47      responses:
48        '201':
49          description: Null response
50        default:
51          description: unexpected error
52          content:
53            application/json:
54              schema:
55                $ref: "#/components/schemas/Error"
```

```
56  /pets/{petId}:
57    get:
58      summary: Info for a specific pet
59      operationId: showPetById
60      tags:
61        - pets
62      parameters:
63        - name: petId
64          in: path
65          required: true
66          description: The id of the pet to retrieve
67          schema:
68            type: string
69      responses:
70        '200':
71          description: Expected response to a valid request
72          content:
73            application/json:
74              schema:
75                $ref: "#/components/schemas/Pet"
76        default:
77          description: unexpected error
78          content:
79            application/json:
80              schema:
81                $ref: "#/components/schemas/Error"
82  components:
83    schemas:
84      Pet:
85        type: object
86        required:
87          - id
88          - name
89        properties:
90          id:
91            type: integer
92            format: int64
93          name:
94            type: string
95          tag:
96            type: string
97      Pets:
98        type: array
99        items:
100          $ref: "#/components/schemas/Pet"
101      Error:
102        type: object
103        required:
104          - code
105          - message
106        properties:
107          code:
108            type: integer
109            format: int32
110          message:
111            type: string
```

# Application state

- Problem: how does a client know that the path for adding a vote is `/poll/{pid}/vote`?
- REST has no standard service contract
- HATEOAS to the rescue

→ GET `/poll/12345`

← 200 OK

Content-Type: `application/xml`

...

`<poll>`

...

`<link rel="vote" href="/poll/12345/vote" />`

...

`</poll>`

# RMMML3 interfaces

```
1  {
2    "_links": {
3      "self": { "href": "/orders/523" },
4      "warehouse": { "href": "/warehouse/56" },
5      "invoice": { "href": "/invoices/873" },
6      "redo": { "href": "/invoices/873/redo" }
7    },
8    "currency": "USD",
9    "status": "shipped",
10   "total": 10.20
11 }
```

# RMML3 interfaces

```
1  {  
2    "_links": {  
3      "self": { "href": "/orders/523" },  
4      "warehouse": { "href": "/warehouse/56" },  
5      "invoice": { "href": "/invoices/873" },  
6      "redo": { "href": "/invoices/873/redo" }  
7    },  
8    "currency": "USD",  
9    "status": "shipped",  
10   "total": 10.20  
11 }
```

The contracts in RMML3 are *links' names*

# RMML3 interfaces

- **Ontologies are the interfaces** in RMML3 REST APIs
- They can be OWL ontologies, JSON-LD vocabularies, folksonomies, it does not matter. But these are your interfaces.



# REST and SOA

- The promise of a lightweight approach to (web) services with diminished burden for the clients holds only within the framework of a specific business model:
  - Large asymmetry between providers and consumers
  - Best effort is a good-enough QoS guarantee

# REST and SOA

- Request/reply model only, no asynchronous calls, no server-initiated messages, no publish/subscribe
- No application level encryption
- No reliable messaging (which plays badly with non-idempotent requests)

# REST and SOA




- No distributed transactions
  - See, for example, “A TS-Based 2PC for Web Services Using Rest Architectural Style” by Luiz A. Hiane S. Maciel and Celso M. Hirata or “Atomic distributed transactions: a RESTful design” by Guy Pardon and Cesare Pautasso.
- No concurrency control
- No contract (no service description language)
  - WSDL 2.0, WADL, OAS, RAML, ...

# REST and SOA

- When QoS is relevant, REST is essentially moving concerns from servers/infrastructure to clients
- More burden for clients developers
- BUT: more awareness too, think, for example, about *graceful degradation*

# Request/reply model only

- Issues
  - HTTP is client/server and client-initiated
  - To receive messages/events/notifications you have to be (1) addressable and (2) able to receive (and process) HTTP requests
- Solutions/workarounds
  - REST (Web) Hooks (requires 1 and 2)
  - Long polling (no requirements)
  - Server-Sent Events tunneling (hackish)
  - Notice that WebSockets are not RESTful



# Long polling

- Issues with polling
  - Short time between retries
    - Network and server overload
  - Long time between retries
    - Clients are not notified promptly
- Long polling
  - When the resource is not available, do not immediately return but keep the HTTP connection open until the resource is available (or until a timeout, in which case the client should retry).
  - See RFC 6202 for known issues

# HTTP methods properties

	SAFE	IDEMPOTENT
GET	YES	YES
PUT	NO	YES
POST	NO	NO
DELETE	NO	YES



# POST Once Exactly

- Use GET to retrieve the URL of a one-shot endpoint
- POST to the endpoint
- In case of failure retry the POST
  - 202 – Accepted  
it had not been processed before
  - 405 – Operation Not Permitted  
the endpoint has been used already

# POST/PUT creation

- A more intuitive (and safer) way of dealing with POST failures when creating resources
- POST /factory
  - Returns the URI of a new empty resource  
/resource/xyz
- PUT /resource/xyz
  - Initializes the resources - this includes whichever state-changing operation on the server side

# Slow creation

→ POST /job

← 202 Accepted

Location: /job/42

→ GET /job/42

← 200 OK

or

← 303 See Other

Location: /job/42/output

# Optimistic concurrency control

- Add €100 to your bank account
  - GET /account/xyz/amount → amount
  - newAmount = amount + 100
  - PUT newAmount → /account/xyz/amount
- What if, in the meanwhile, somebody else tries the same?

# Optimistic concurrency control

Use ETags and HTTP conditional requests

GET /account/xyz/amount

200 – OK

ETag: qwerty

PUT /account/xyz/amount

If-Match: qwerty

# Try-Confirm/Cancel

Distributed transactions the REST way

- Create explicit representations of to-be-confirmed operations as resources (**Try**)
- **Confirm** by PUTting on these resources (PUT is idempotent, you can use retry)
- **Cancel** is performed by the services themselves after a timeout
- GET on an intermediate resource returns time left before cancel

# Binary messaging

- JSON is designed to be human-readable
- JSON is the lingua franca in the API world
- JSON parsing can be the largest workload of a REST API

# Binary alternatives

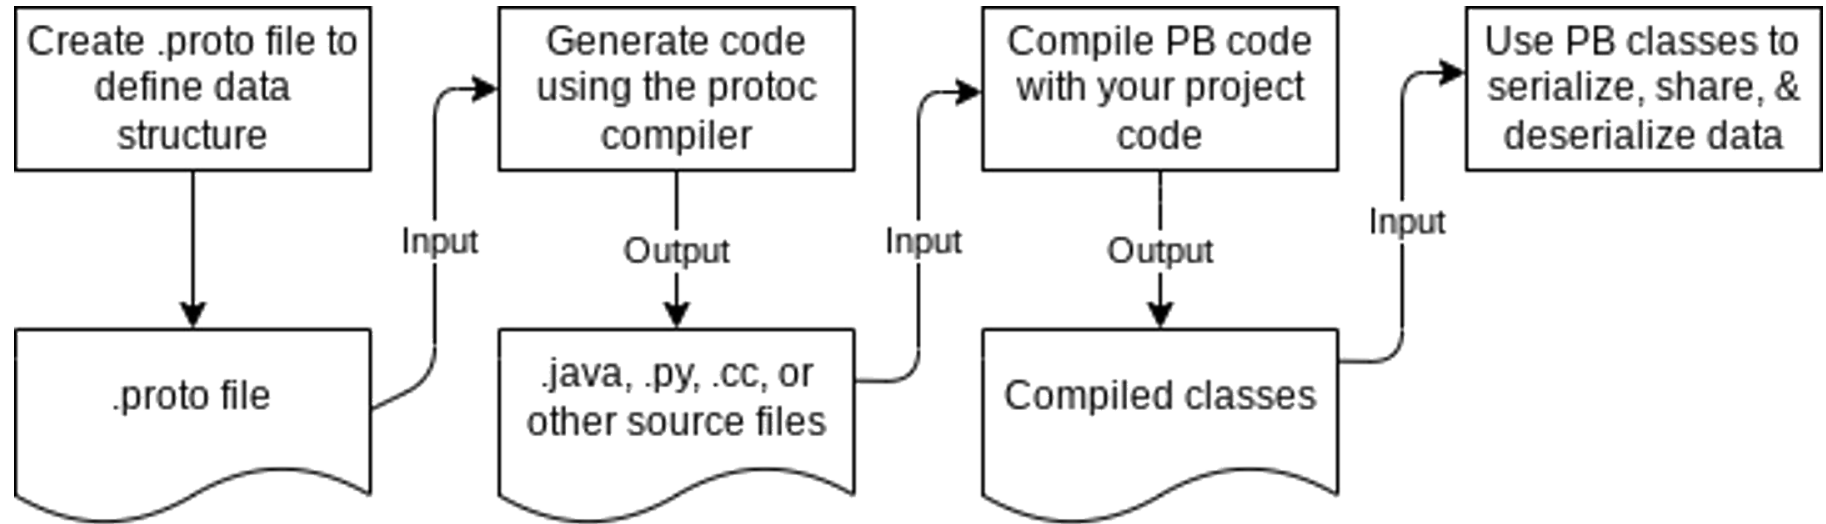
- Several binary formats have been proposed, only a few are designed to really promote interoperability
- Usual approach
  - Define a schema with a type definition language
  - Map types to specific programming languages structures



# Binary serialization

- CORBA (1991 - serialization is just a part of the whole standard)
- Apache Thrift (originally a FB project)
- Protocol Buffers (originally a Google project)
- Apache Avro (from the Hadoop project)
- ...

# Protocol buffers workflow



# .proto file example

```
message Person {  
    optional string name = 1;  
    optional int32 id = 2;  
    optional string email = 3;  
}
```

# RPC vs binary serialization

- Binary serialization does not imply the use of a specific communication infrastructure
- You can do REST using a binary format to transfer resources
- Some binary serialization systems include an RPC stack (Thrift) or have a preferred RPC stack (like gRPC and Protocol Buffers)

# gRPC

- gRPC is an RPC framework that uses protocol buffers to specify and (de)serialize data
- It's implemented on top of HTTP/2
- Its main distinctive feature is the support for **bi-directional streaming**
- It directly supports C#, C++, Dart, Go, Java, Kotlin, Node, Objective-C, PHP, Python and Ruby

# gRPC

